

Java Multithreading, Concurrency & Performance Optimization

- Udemmy <https://www.udemy.com/course/java-multithreading-concurrency-performance-optimization/learn/lecture/10187964#overview>
- Github <https://github.com/jgregorio0/java-multithreading>

Introduction

Why to use multithreading?

Responsiveness / Concurrency

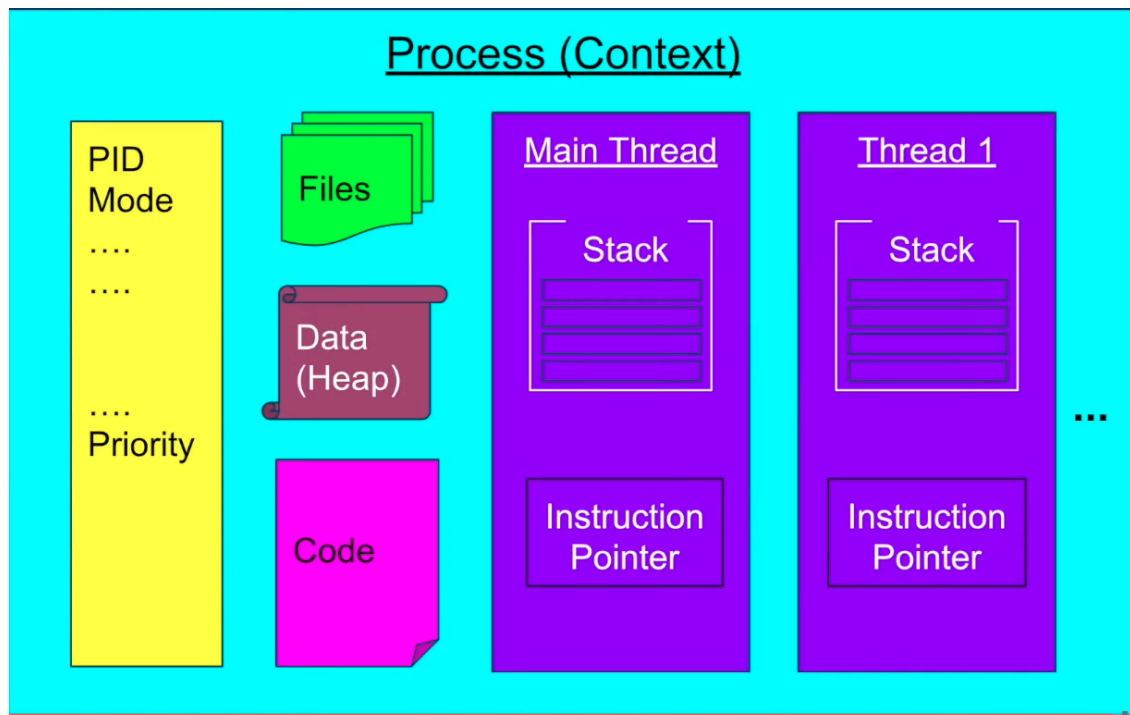
- Serve multiple users simultaneously
- Critical in user interface, ie: video player
- One core creates illusion of multiple tasks in parallel

Performance / Parallelism

- Multiple cores run tasks in parallel
- Higher performance >> more work in same time and less machines

Threads structure in OS

1. Users run an app >> OS create an instance (process / context) of the app from HD to memory which is completely isolated from the other processes
2. Each **process** contains
 - a. metadata (PID, priority, mode,...)
 - b. files
 - c. code
 - d. heap (data for our app)
 - e. One (Main Thread) or more threads
3. **Thread** contains
 - a. Stack: local variables are stored
 - b. Instruction pointer: address of the next instruction
4. All but the stack and instruction pointer is shared



Context switch

- When switching between threads we need to
 - stop a current thread execution
 - schedule it out
 - schedule new thread in
 - start new thread

Take into account that:

- **Thrashing** happens when
 - management thread > productive work
- Threads resources < process resources
 - The costs of changing thread from the same resources < threads from diff process

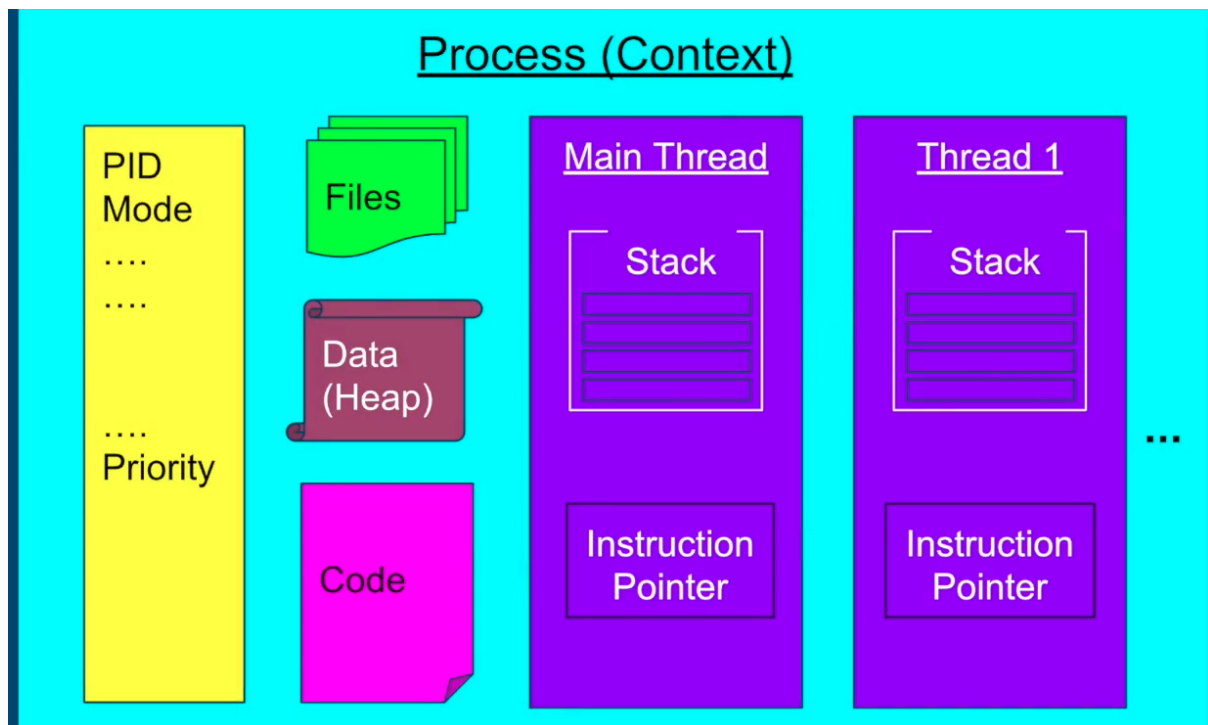
Thread scheduling

- Who run first?
 - Long threads can cause **starvation** over short threads
 - Short threads can cause long threads are never executed
 - Solving starvation by dividing time into **Epochs** and adding bonus to threads did not complete in the previous epoch
 - Epoch (time slice) in which a thread uses CPU
 - dynamic priority = static priority + Bonus
 - static priority
 - UI will be executed first improving UX
 - Bonus

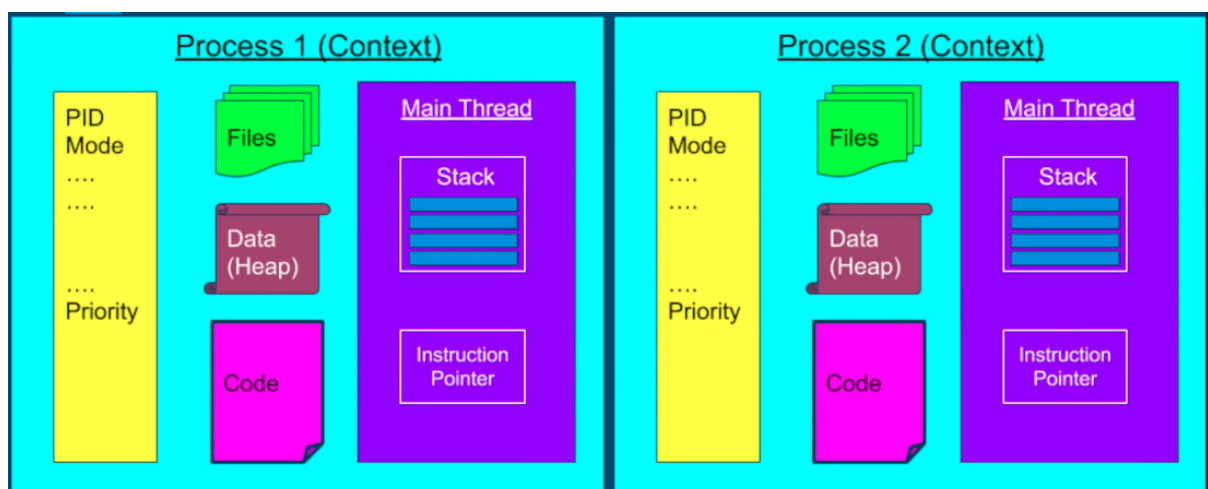
- Increase priority for threads that did not complete in last epoch prevents starvation

Multithreaded vs Multi Processes

- Multi Threads
 - share a lot of data
 - Threads are faster to create and destroy
 - Shorter context switches



- Multi Processes
 - Security and stability is higher
 - Task are unrelated



Thread fundamentals

Create thread implementing Runnable

```
public class Main {
    public static void main(String[] args) {
        Thread t = new Thread(() -> {
            System.out.println("We are in thread " +
Thread.currentThread().getName());
        });
        t.start();
    }
}
```

- setName
- setPriority
 - Thread.MAX_PRIORITY
- Thread.sleep(milliseconds)
- setUncaughtExceptionHandler

Create thread extending Thread

```
public class SimpleThread2 {
    public static void main(String[] args) throws InterruptedException {
        Thread rt = new RunnableThread();
        rt.start();
    }

    static class RunnableThread extends Thread {
        @Override
        public void run() {
            System.out.println("We are in thread " +
Thread.currentThread().getName());
        }
    }
}
```

Interrupt thread

- Threads consume resources
 - Memory
 - kernel resources
 - CPU
 - cache
- Thread is finished but app is still running >> Clean up threads' resources
- Thread is misbehaving
- App will not stop if one thread is still running >> close all threads gracefully

InterruptedException

- InterruptedException: Thread executes a method that throws InterruptedException

```
private static class BlockingTask implements Runnable {
    @Override
    public void run() {
        try {
            Thread.sleep(50000);
        } catch (InterruptedException e) {
            System.out.println("Exiting blocking task");
        }
    }
}
```

isInterrupted

- isInterrupted: Theads code is handling interrupt signal

```
@Override
public void run() {
    System.out.println(base + "^" + power + " = " + pow(base, power));
}

private BigInteger pow(BigInteger base, BigInteger power) {
    BigInteger result = BigInteger.ONE;
    for (BigInteger i = BigInteger.ZERO; i.compareTo(power) != 0; i =
i.add(BigInteger.ONE)) {

        if (Thread.currentThread().isInterrupted()) {

            System.out.println("Prematurely interrupted computation");
            return BigInteger.ZERO;
        }
        result = result.multiply(base);
    }
    return result;
}
```

Daemon threads

- Background tasks that should not block our app
 - File saving
- Code could not listen to interrupt
 - External library

setDaemon

- **setDaemon(true)**

- Thread will be **finished** even if not throwing **InterruptedException** or checking **isInterrupted**

Thread coordination

Join

- join method sleep thread A until thread B is finished
 - timeout parameter throws InterruptedException after x milliseconds
-

Performance optimization

- Latency: Time to complete a task. Time units
- Throughput: Amount of task in a period. Task/time units

Latency

- Latency = T/N
 - T time to execute original task
 - N number of subtasks
- Dividing a task into N tasks to run in parallel
 - N = number of cores
 - each core running 1 thread
 - if threads are runnable without interruption (IO blocking, sleep...)
 - No other tasks are consuming CPU
 - Hyperthreading: virtual cores share hardware
- Cost of parallelization
 - Breaking task into multiple tasks
 - Thread creation and passing task to threads
 - Time to schedule a task
 - Time until last task finishes and signals
 - Time until aggregator task runs
 - Aggregation of the results
- It is not possible to divide a task always:
 - Parallelizable tasks
 - Sequential tasks
 - Partially parallelizable / partially sequential

Throughput

- Throughput = N/T (rendimiento)
 - N number of subtasks
 - T time to execute original task
- Dividing tasks into N tasks
 - Latency = T/N
- Running tasks in parallel

- Each task in different thread
- Improve throughput by N
- $N = \text{Threads} = \text{Cores}$
- Thread pooling
 - Reusing threads minimize creation and schedule tasks
 - `Executor.newFixedThreadPool`

JMeter

Automate test performance

1. Add Thread Group
 - a. 200 threads
2. Add Logic Controller / While Controller
3. Add Config Element / CSV Data Set Config
 - a. Filename
 - i. search_words.csv
 - b. Variable Names
 - i. WORD
 - c. Delimiter
 - i. \n
 - d. Recycle on EOF
 - i. false
 - e. Stop thread on EOF
 - i. true
4. While Controller Condition
 - a. `${__javaScript("${WORD}" !== "<EOF>")}`
5. Add Sample / HTTP Request
 - a. protocol
 - i. http
 - b. server
 - i. localhost
 - c. port
 - i. 8000
 - d. Method
 - i. GET
 - e. Path
 - i. /search?word=\${WORD}
6. Add listener / Summary Report
7. Add listener / View Results Tree

Thread Data Sharing

Stack


- Primitive types declared locally will be stored on the stack only
- Other threads have no access to the stack

- The following example shows how
 - First stack frame contains x and y
 - second stack frame contains a, b and s (return s)
 - second stack is invalidated and its result is allocated in first frame
 - First stack frame contains result of second frame and its invalidated when finished

What is the stack?

```
void main(String [] args) {  
    int x = 1;  
    int y = 2;  
    int result = sum(x, y) ;  
}
```

```
int sum(int a, int b) {  
    int s = a + b;  
    return s;  
}
```




```

1 public class Stack {
2     public static void main(String[] args) { args: {}
3         int x = 1; x: 1
4         int y = 2; y: 2
5         int res = sum(a: 1, b: 2);
6     }
7
8     private static int sum(int a, int b) {
9         int s = a + b;
10        return s;
11    }
12 }

```

Debug: Stack x

Debugger Console

Frames

- main:5, Stack

Variables

- args = {String[0]@468} []
- x = 1
- y = 2

```

1 public class Stack {
2     public static void main(String[] args) {
3         int x = 1;
4         int y = 2;
5         int res = sum(a: 1, b: 2);
6     }
7
8     private static int sum(int a, int b) { b: 2
9         int s = a + b; b: 2 s: 3
10        return s; s: 3
11    }
12 }

```

Debug: Stack x

Debugger Console

Frames

- sum:10, Stack
- main:5, Stack

Variables

- a = 1
- b = 2
- s = 3

```

1 public class Stack {
2     public static void main(String[] args) { args: {}
3         int x = 1; x: 1
4         int y = 2; y: 2
5         int res = sum(a: 1, b: 2); res: 3
6     }
7
8     private static int sum(int a, int b) {
9         int s = a + b;
10        return s;
11    }
12 }

```

Debug: Stack x

Debugger Console

Frames

- main:6, Stack

Variables

- args = {String[0]@468} []
- x = 1
- y = 2
- res = 3

Heap

- Static variables, objects and primitive types that objects contains will be stored on the heap
- Heap is shared between threads
- Governed by Garbage Collector
 - It should remove objects when there are no references to the objects
 - Static variables stay forever

References

- Reference is a pointer to an Object
- Local reference is allocated in the Stack
- but If they are member of a class they will be allocated in the Heap

Concurrency Solutions

Atomic operations

- Occurred at once
- Single state, all or nothing, without intermediate states
- What is atomic?
 - Assignment to
 - References, including getters and setters
 - Assignment to
 - int
 - short
 - byte
 - float
 - char
 - boolean
 - ~~long (64 bits)~~
 - ~~double (64bits)~~
 - volatile long
 - volatile double
 - java.util.concurrent.atomic

Synchronized

- Lock mechanism
- Locks all synchronized methods of the object
- Use sychronized blocks instead of synchronized methods lock only the block instead of all methods

Data Race

- A shared resource
 - is accessed by multiple threads
 - is modified by at least one of those thread
- Timing of threads scheduling may cause incorrect results
- **Problem:** Non atomic operations performed
- **Solution:**
 - Using **synchronized** to atomize a critical section
 - Using **volatile** to atomize long and double assignaments
- **Example:**

Thread1 { x++; y++;}

Thread2 { if (y > x) throws new Exception("Data Race detected!!"); }

Race condition

- A shared resource
 - is accessed by multiple threads
 - is modified by at least one of those thread
- Compiler & CPU may execute code out of order to increase performance and utilization maintaining logical correctness
- **Problem:** Reorder code in one thread results in unexpected behaviour for the other
- **Solution:**
 - Using **synchronized** to atomize a critical section
 - Using **volatile** to guarantee order on the previous and next instruction
- **Example:**

Thread1 { i++ }

Thread2 { i-- }

Result i != 0

Summary

- **synchronized**
 - **atomize** a critical section
 - performance decrease
- **volatile**
 - **atomize** long and double assignments
 - **guarantee order** on the previous and next instruction
- **Rule**
 - Any **variable** used by **multiple threads** and **modified by one** at least must be in **synchronized** block **or** be declared as **volatile**

Lock

- Coarse grain locking
 - one lock for all shared resources
 - Decrease paralelism
- Fine grain locking
 - many locks for each shared resource
 - Deadlock

Deadlock

| Step | Thread 1 | Thread 2 |
|------|----------|----------|
| 1 | lock A | |
| 2 | | lock B |

| | | |
|----------|------------------|------------------|
| 3 | lock B > BLOCKED | |
| 4 | | lock A > BLOCKED |
| DEADLOCK | | |

Deadlock condition

- **Mutual Exclusion**
 - Only one thread can have **exclusive access to a resource**
- **Hold and Wait** - At least one thread is holding a resource and is waiting for another resource
- **Non-preemptive allocation** - A resource is released only after the thread is done using it.
- **Circular wait** - A chain of at least two threads each one is holding one resource and waiting for another resource

Deadlock solution

- Avoid circular wait
 - Strict **order of locking** shared resources

| Step | Thread 1 | Thread 2 |
|------|----------|---------------|
| 1 | lock A | |
| 2 | | lock A > WAIT |
| 3 | lock B | |
| 4 | unlock B | |
| 5 | unlock A | |
| 6 | | lock A |
| 7 | | lock B |
| 8 | | unlock B |
| | | unlock A |

- Deadlock detection - Whatchdog
- Thread interruption
- tryLock operations

ReentrantLock

- same functionalities than synchronized method
- plus:
 - check lock status
 - lockInterruptibly
 - Allow to **interrupt** thread **waiting** for lock
 - tryLock
 - thread is not suspended forever
 - **wait** until **timeout** is achieved and return false if !lock

ReentrantReadWriteLock

- same than ReentrantLock but
 - it allow **many** threads to **read**
 - only **one write**.
 - It also block readers when writing

Semaphore

- Restrict how many threads access to shared resources
- **Lock**
 - allows access to **single thread** only
 - is **reentrant**
- **Semaphore**
 - allows access to a **multiple threads**
 - **Not reentrant**

Inter-thread communication

Condition Variables

- Thread can wait for condition that release other thread using Lock.newCondition()
- java.util.concurrent.locks.Lock
 - lock
 - unlock
- java.util.concurrent.locks.Condition
 - await
 - signal
 - signalAll

Object

- java.lang.Object
 - synchronized(object)
 - wait

- notify
- notifyAll
- current thread waits until other thread wakes it up
- notify
 - wakes up 1 thread waiting on that object
- notifyAll
 - wakes up all threads waiting on that object

Lock-Free Algorithms, Data-Structures & Techniques

Atomic Objects

- PROS
 - simplicity
 - No need for locks or synchronization
 - No race conditions or data races
- CONS
 - Only the operation is atomic
 - Race condition between 2 separate atomic operations could appear

AtomicReference

- Wraps a reference and allow to perform atomic operations on the reference
- if currentValue == expectedValue > assign newValue
- else if currentValue != expectedValue > nothing